

gem5 tutorial

References:

- <https://github.com/gem5bootcamp>



Environment

The steps outlined in this tutorial have been tested on the computers located in **Delta Building, Room 219**.

These computers are equipped with:

- CPU: 13th Gen Intel(R) Core(TM) i5-13500
- RAM: 16 GB
- SSD: 954 GB

Docker Desktop, WSL, Hyper-V, KVM are pre-installed.

Installation on Your Own Computer

You may also install Docker on your own computer. Please ensure you have at least **10 GB** of disk space available.

This instructions for **Windows** is provided.



Install WSL2 on Windows

STEP 1: Enable Windows Features

1. Open **File Explorer** and type Control **Panel**(控制台)\Programs(程式集) into the address bar, then press Enter.
2. Click on **Turn Windows features on or off**(開啟或關閉Windows功能).
3. Check the boxes for **Windows Subsystem for Linux** and **Virtual Machine Platform**.

If these options are missing, you may need to enable **virtualization technology** (Intel VT-x / AMD-V) in your computer's BIOS/UEFI settings.

4. Click OK, then select Restart now to reboot your computer.

STEP 2: Install WSL2 using PowerShell

1. Open **PowerShell** as Administrator.
2. Run the following command to install WSL2 and the default Linux distribution (Ubuntu):

```
ws1 --update --web-download
```

3. Wait for the installation to complete. Once finished, restart your computer if prompted.

Install Docker Desktop on Windows

For Windows users:

1. Download [the installer](#).
2. Double-click `Docker Desktop Installer.exe` to run the installer.
3. Follow the instructions in the installation wizard to authorize the installer and proceed with the installation.
4. Once the installation is complete, select **Close** to finish the process.

Back up files from `$HOME\workspaces\` to a flash drive

Delta 219 computers are configured to automatically revert to a previously saved system state (a restore point) without prior notice.

- Your files may be permanently deleted without warning at any time (unannounced data loss).
- You cannot predict when these reverts will occur, nor can you prevent them from happening (limited user control).

Temporary data storage:

Your files are temporarily stored in the `$HOME\workspaces\` directory. These files are deleted during the automatic system revert process.

Protecting your data:

To avoid data loss, regularly back up your data using external USB drives or cloud storage services.

Reason for automatic system reverts:

This feature is designed to maintain a consistent simulation environment for all users.



Open Docker Desktop on Windows

For Windows users:

1. Left-click on the **Start button** (or press the `Windows key`).
2. Select **Docker Desktop** (or type `Docker Desktop` and press `Enter`).
3. Wait until Docker Desktop has fully started

Expected outcome: Docker Desktop is open and fully initialized, ready to run containers.



Open a terminal emulator on Windows

For Windows users:

1. Right-click on the **Start button** (or press the `Windows key` + `X`).
2. Select **Windows PowerShell** (or press `I`).

Expected outcome: A terminal emulator will open.

```
Windows PowerShell
```

```
Copyright (C) Microsoft Corporation. 著作權所有，並保留一切權利。
```

```
請嘗試新的跨平台 PowerShell https://aka.ms/pscore6
```

```
PS C:\Users\user>
```

Run a container in PowerShell

For Windows Pro / Enterprise:

If you are on **Windows Pro/Enterprise**, enable **Hyper-V** and **Nested Virtualization** so that **KVM** can be passed through into the container. (**Computers in 219 are available for use**)

```
docker run -it `
--device /dev/kvm `
--name ee6455-gem5 `
--volume $HOME\workspaces\:./workspaces/ `
--workdir /workspaces/2025/ `
--hostname EE6455-gem5 `
ghcr.io/gem5/devcontainer:bootcamp-2024
```

For Windows Home:

If you are using **Windows Home edition**, **Hyper-V** and **nested virtualization** are not available.

In this case, **do not add** `--device /dev/kvm`, since **KVM** passthrough will not work.

```
docker run -it `
--name ee6455-gem5 `
--volume $HOME\workspaces\:./workspaces/ `
--workdir /workspaces/2025/ `
--hostname EE6455-gem5 `
ghcr.io/gem5/devcontainer:bootcamp-2024
```

Expected outcome: An interactive TTY will indicate its readiness to accept commands.

```
root@EE6455-gem5: /workspaces/2025#
```

Alternative outcome: A newer image will be downloaded before the interactive TTY displays the username, hostname, and working directory.

```
PS C:\Users\user> docker container run -it `
>> --name ee6455-gem5 `
>> --volume $HOME\workspaces\: /workspaces/ `
>> --workdir /workspaces/2025/ `
>> --hostname EE6455-gem5 `
>> ghcr.io/gem5/devcontainer:v25-0
Unable to find image 'ghcr.io/gem5/devcontainer:v25-0' locally
v25-0: Pulling from gem5/devcontainer
00d679a470c4: Pull complete
c782a11a41b6: Pull complete
4f6b9996da3d: Pull complete
bb60cbcef558: Pull complete
92c0a1bbbf0cc: Pull complete
```

Note on KVM and Full-System Simulation:

In several chapters of the gem5 bootcamp, **KVM** acceleration is used to speed up simulation.

- **Full-System (02-Using-gem5/07-full-system)** and **CHI Protocol (03-Developing-gem5-models/07-chi-protocol)** – these chapters use **KVM acceleration** to speed up full-system simulations (e.g., booting Ubuntu or running workloads with advanced systems).

Workload Example (07-full-system):

- `x86-ubuntu-24.04-boot-with-systemd`
 - With KVM: ~1–2 minutes
 - Without KVM: 3+ hours

(continued on next page)

Supported Host OS:

- **Windows Pro / Enterprise with Hyper-V enabled**

If **nested virtualization** is enabled and you run the container with `--device /dev/kvm`, then KVM acceleration will work.

See Microsoft's official docs:

- [Overview of Nested Virtualization](#)
- [Enable Nested Virtualization](#)

- **Linux Host OS**

If your CPU/BIOS supports virtualization (Intel VT-x or AMD-V) and `/dev/kvm` is available, you can pass it into the container with `--device /dev/kvm`.

Not Supported Host OS:

- **Windows Home Edition**

Hyper-V and nested virtualization are not supported. In this case, KVM cannot be used and simulations will run **without hardware acceleration**, which will be significantly slower.



- **macOS (Intel & Apple Silicon)**

macOS does not expose `/dev/kvm` to Docker. Even on Apple Silicon (M1/M2/M3/M4), KVM cannot be used inside containers.

Simulations must therefore run **without hardware acceleration**, which will be much slower.

Important: KVM requires the host ISA to match the guest ISA:

- x86 host → can accelerate x86 guest (e.g., Ubuntu x86 FS workloads)
- ARM host → can accelerate ARM guest
- Cross-ISA acceleration (e.g., x86 host running RISC-V guest) is **not supported**.

Manage an Existing Docker Container

- Start an existing container:

```
docker start ee6455-gem5
```

- Attach a shell inside a running container:

```
docker exec -it ee6455-gem5 bash
```

- Opens a new Bash shell inside the container.
- You can have multiple terminals connected at the same time.

- Stop a running container:

```
docker stop ee6455-gem5
```



- Remove a container

```
docker rm ee6455-gem5
```

- Works only if the container is stopped.
- **Warning:** Removing a container will delete all data created inside it (unless you used volumes or bind mounts to persist the data).

Get the source code of gem5

Clone the Bootcamp Repository (inside container):

```
time git clone --recurse-submodules \
https://nas.larc-nthu.net:8443/ee6455_2025/public-gem5bootcamp-2025 /workspaces/2025/
```

Expected outcome: A repository will be cloned, including its submodules.

```
root@EE6455-gem5:/workspaces/2025# time git clone --recurse-submodules https://nas.larc-nthu.net:8443/ee6455_2025/public-gem5bootcamp-2025 /workspaces/2025/
Cloning into '/workspaces/2025'...
warning: redirecting to https://nas.larc-nthu.net:8443/ee6455_2025/public-gem5bootcamp-2025.git/
remote: Enumerating objects: 11106, done.
remote: Counting objects: 100% (130/130), done.
remote: Compressing objects: 100% (130/130), done.
remote: Total 11106 (delta 78), reused 0 (delta 0), pack-reused 10976 (from 1)
Receiving objects: 100% (11106/11106), 385.46 MiB | 11.07 MiB/s, done.
Resolving deltas: 100% (7510/7510), done.
Updating files: 100% (966/966), done.
Submodule 'gem5' (https://github.com/gem5/gem5) registered for path 'gem5'
Submodule 'gem5-resources' (https://github.com/gem5/gem5-resources) registered for path 'gem5-resources'
Cloning into '/workspaces/2025/gem5'...
remote: Enumerating objects: 295271, done.
remote: Counting objects: 100% (604/604), done.
remote: Compressing objects: 100% (329/329), done.
remote: Total 295271 (delta 460), reused 275 (delta 275), pack-reused 294667 (from 3)
Receiving objects: 100% (295271/295271), 171.33 MiB | 5.20 MiB/s, done.
Resolving deltas: 100% (225064/225064), done.
Cloning into '/workspaces/2025/gem5-resources'...
```

Preserve downloaded workloads across container recreations

By default, gem5 downloads resources (kernels, disk images, benchmarks) into `/root/.cache/gem5` inside the container.

If the container is deleted, these downloads will also be lost.

To avoid re-downloading large workloads every time, you can set an environment variable to redirect the resource cache to a persistent directory:

```
export GEM5_RESOURCE_DIR=/workspaces/resources
```

This way, when the container is removed and recreated, previously downloaded workloads will still be available.



Run a simulation using gem5

Run this command inside the container:

```
time gem5-mesi --outdir=/workspaces/m5out/ \
/workspaces/2025/materials/01-Introduction/02-getting-started/completed/basic.py
```

Expected outcome:

- The `X86DemoBoard` will be simulated using `x86-ubuntu-24.04-img` as the workload.
- On first run, gem5 will automatically download required resources:
 - `x86-linux-kernel-5.4.0-105-generic`
 - `x86-ubuntu-24.04-img` (~4.9 GB , takes some time to download and decompress)

Note: Depending on your network speed, downloading and unpacking `x86-ubuntu-24.04-img` may take several minutes.



```

root@EE6455-gem5:/workspaces/2025# time gem5-mesi --outdir=/workspaces/m5out/ \
/workspaces/2025/materials/01-Introduction/02-getting-started/completed/basic.py
gem5 Simulator System. https://www.gem5.org
gem5 is copyrighted software; use the --copyright option for details.

gem5 version 24.0.0.0
gem5 compiled Jul 25 2024 18:47:27
gem5 started Aug 21 2025 09:19:18
gem5 executing on EE6455-gem5, pid 14
command line: gem5-mesi --outdir=/workspaces/m5out/ /workspaces/2025/materials/01-Introduction/02-getting-started/completed/basic.py

warn: The X86DemoBoard is solely for demonstration purposes. This board is not known to be representative of any real-world system. Use with caution.
info: Using default config
Resource 'x86-linux-kernel-5.4.0-105-generic' was not found locally. Downloading to '/root/.cache/gem5/x86-linux-kernel-5.4.0-105-generic'...
Finished downloading resource 'x86-linux-kernel-5.4.0-105-generic'.
Resource 'x86-ubuntu-24.04-img' was not found locally. Downloading to '/root/.cache/gem5/x86-ubuntu-24.04-img.gz'...
Finished downloading resource 'x86-ubuntu-24.04-img'.
Decompressing resource 'x86-ubuntu-24.04-img' ('/root/.cache/gem5/x86-ubuntu-24.04-img.gz')...
Finished decompressing resource 'x86-ubuntu-24.04-img'.
warn: Max ticks has already been set prior to setting it through the run call. In these cases the max ticks set through the `run` function is used
Global frequency set at 100000000000 ticks per second
src/mem/dram_interface.cc:690: warn: DRAM device capacity (8192 Mbytes) does not match the address range assigned (2048 Mbytes)
src/sim/kernel_workload.cc:46: info: kernel located at: /root/.cache/gem5/x86-linux-kernel-5.4.0-105-generic
src/base/statistics.hh:279: warn: One of the stats is a legacy stat. Legacy stat is a stat that does not belong to any statistics::Group. Legacy stat is deprecated.
    0: board.pc.south_bridge.cmos rtc: Real-time clock set to Sun Jan  1 00:00:00 2012
board.pc.com_1.device: Listening for connections on port 3456
src/base/statistics.hh:279: warn: One of the stats is a legacy stat. Legacy stat is a stat that does not belong to any statistics::Group. Legacy stat is deprecated.
src/dev/intel_8254_timer.cc:128: warn: Reading current count from inactive timer.
board.remote_gdb: Listening for connections on port 7000
src/sim/simulate.cc:199: info: Entering event queue @ 0. Starting simulation...
src/mem/ruby/system/Sequencer.cc:680: warn: Replacement policy updates recently became the responsibility of SLICC state machines. Make sure to setMRU() near callbacks in .sm files!
build/ALL/arch/x86/generated/exec-ns.cc.inc:27: warn: instruction 'fninit' unimplemented

real    42m27.091s
user    0m38.201s
sys     0m16.827s
root@EE6455-gem5:/workspaces/2025#

```

Retrieve the simulation statistics

- On Windows

After the simulation finishes, open `$HOME\workspaces\m5out\stats.txt` using a text editor. This file contains the detailed statistics generated by gem5.

- Inside the container

You can also view the file directly inside the running container, for example by scrolling through the whole file:

```
root@EE6455-gem5:/workspaces/2025# less /workspaces/m5out/stats.txt
```

Expected outcome: The first few lines will resemble this:

```
----- Begin Simulation Statistics -----
simSeconds          0.020000          # Number of seconds simulated (Second)
simTicks            20000000000        # Number of ticks simulated (Tick)
finalTick           20000000000        # Number of ticks from beginning of simulation (restored from checkpoints and never reset) (Tick)
simFreq             1000000000000       # The number of ticks per simulated second ((Tick/Second))
hostSeconds         18.88              # Real time elapsed on the host (Second)
hostTickRate        1059100397         # The number of ticks simulated per host second (ticks/s) ((Tick/Second))
hostMemory          2770924            # Number of bytes of host memory used (Byte)
simInsts            7479814            # Number of instructions simulated (Count)
simOps              34912342           # Number of ops (including micro ops) simulated (Count)
hostInstRate        396059             # Simulator instruction rate (inst/s) ((Count/Second))
hostOpRate          1848597            # Simulator op (including micro ops) rate (op/s) ((Count/Second))
board.cache_hierarchy.ruby_system.delayHistogram::bucket_size 2          # delay histogram for all message (Unspecified)
board.cache_hierarchy.ruby_system.delayHistogram::max_bucket 19          # delay histogram for all message (Unspecified)
board.cache_hierarchy.ruby_system.delayHistogram::samples 735551          # delay histogram for all message (Unspecified)
board.cache_hierarchy.ruby_system.delayHistogram::mean 1.036855          # delay histogram for all message (Unspecified)
board.cache_hierarchy.ruby_system.delayHistogram::stdev 2.687016          # delay histogram for all message (Unspecified)
```

Retrieve the simulation configuration

- On Windows: open `$HOME\workspaces\m5out\config.ini` using a text editor.
- Inside the container:

```
less /workspaces/m5out/config.ini
```

Expected outcome: The first few lines will resemble this:

```
[board]
type=System
children=cache_hierarchy clk_domain dvfs_handler iobus memory pc processor workload
auto_unlink_shared_backstore=false
cache_line_size=64
eventq_index=0
exit_on_work_items=true
init_param=0
m5ops_base=4294901760
```

```
mem_mode=timing
mem_ranges=0:2147483648 3221225472:3222274048
memories=board.memory.mem_ctrl.dram
mmap_using_noreserve=false
multi_thread=false
num_work_ids=16
readfile=
redirect_paths=
shadow_rom_ranges=
shared_backstore=
symbolfile=
thermal_components=
thermal_model=NULL
work_begin_ckpt_count=0
work_begin_cpu_id_exit=-1
work_begin_exit_count=0
work_cpus_ckpt_count=0
work_end_ckpt_count=0
work_end_exit_count=0
work_exit_id=1
```